

# Tutorial

## Introduction to MPI on python using mpi4py



[https://github.com/jeudy/tutorial\\_mpi4py](https://github.com/jeudy/tutorial_mpi4py)

# Tutorial requisites

- Python 2.7+
- Numpy
  - `pip install numpy`
- An implementation of MPI (openmpi or mpich)
  - `sudo apt-get install openmpi-bin openmpi-common libopenmpi-dev`
- Mpi4py
  - `pip install mpi4py`



SciPyLA  
2016

# Parallel programming

- Distribute a program's work among several execution threads or processors.
- Today's computers come with multiple processors/cores.
- A typical program executes its instructions sequentially using a single processor
- Each processor has its own private memory space.


# Parallel programming

- **OpenMP**: Takes advantage of the multi core architecture capable of executing several threads in parallel: **Shared memory**.
- **MPI**: several computers connected using the network forming a distributed system (**distributed memory**), potentially having thousands of processors available for processing.



SciPyLA  
2016

# Beowulf Cluster



- Message Passing Interface
- Designed in 1992 in a supercomputing conference to decide a standard.
- Its a model where a program passes messages between processes to execute a task.
- Example: a central process assigns work to "slave" processes, transmitting the data to use.

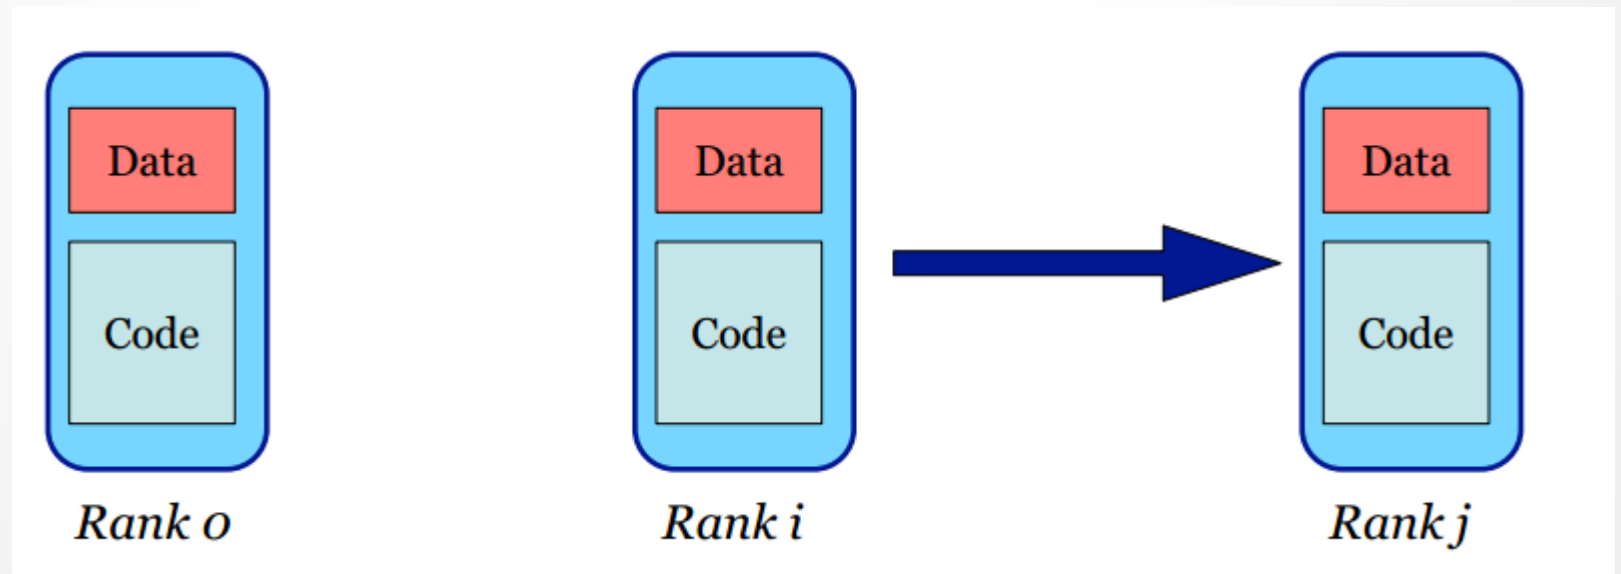
- Its a **library** that has special functions to allow communication and synchronization between processes.
- Implemented in many programming languages. Has wrappers for C and Fortran compilers (mpicc and mpif90).

- Open-source implementations:
  - Open Mpi: <http://www.open-mpi.org/>
  - Mpich: <http://www.mpich.org/>
  - **Mpi4py**: uses some of the previous implementations, from python, without any additional compilers (<http://mpi4py.readthedocs.io/>).
- MPI uses a special launcher to run programs:  
mpirun:
  - **mpirun** -np 10 program.exe



# MPI Concepts

- **Communicator:** group of processes that can communicate to each other.
- **Rank:** numeric identifier for each process within a communicator. Values usually go from 0 to N-1



- Created by Lisandro Dalcin.
- Uses an object oriented approach, keeping the standard semantics of MPI.
- Can communicate any native python object, or user defined objects using the **pickle** module for serialization.
- To run a parallel python program:

```
mpirun python -np 10 myprogram.py
```

# Initialization

- In general, MPI programs need to initialize and release the environment through function calls: Init and Finalize.
- In mpi4py case, the Init call is not critical. The environment is automatically initialized when the library is imported.

```
import mpi4py.MPI as MPI

if not MPI.Is_initialized():
    MPI.Init()

MPI.Finalize()
```

- From the program, you can access the Communicator information, and each process can get its **rank**.
  - COMM\_WORLD represents the Communicator
  - .Get\_size()
  - .Get\_rank()



**SciPyLA**  
**2016**

# MPI's Hello World

```
"""
Simplest mpi program
To execute this run: mpirun -np 4 python hello.py
"""

import mpi4py.MPI as MPI

if not MPI.Is_initialized():
    MPI.Init()

comm = MPI.COMM_WORLD
myid = comm.Get_rank()
size = comm.Get_size()

print "I am process {id}. Total size: {size}\n".format(id=myid, size=size)

MPI.Finalize()
```

**hello.py**

# MPI's Hello World

- `mpirun -np 10 python hello.py`

```
I am process 5. Total size: 10
I am process 0. Total size: 10

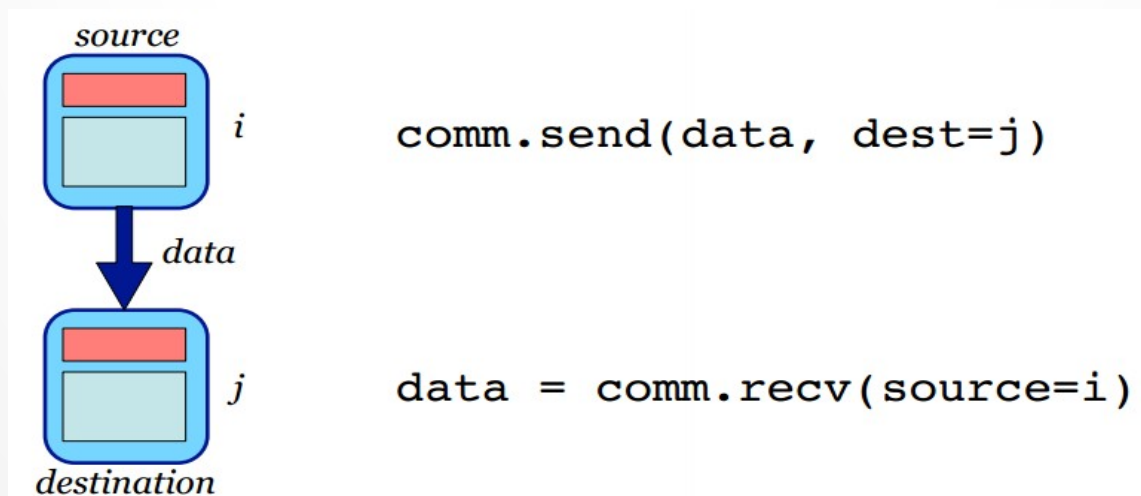
I am process 2. Total size: 10
I am process 3. Total size: 10
I am process 6. Total size: 10
I am process 7. Total size: 10

I am process 8. Total size: 10
I am process 4. Total size: 10

I am process 1. Total size: 10
I am process 9. Total size: 10
```

# Point to Point communication

- A process can send messages to another process directly, using its rank as id.
- Two operations are defined: **send** and **receive**.
- May be synchronous or asynchronous.



# Point to Point communication

- **send** y **recv** transmit native or user created python objects using pickle. Synchronous (blocking).
- `isend` and `irecv` are asynchronous. They return *Request* objects with methods `test()` and `wait()`
- **Send** y **Recv** transmit buffers (i.e: numpy arrays)
- Same pattern on most functions (lowercase for *python objects*, uppercase for *buffers*).



# Synchronous send / recv sample

- `send_recv_object.py`
- `send_recv_buffer.py`

# Asynchronous send / recv sample

- `send_recv_async.py`
- `send_recv_async_wait.py`



*SciPyLA*  
2016

## Example: Ping Pong

- A message is exchanged between 2 processes, 10 times, carrying a counter.
- The process with rank 1 increments the counter when it receives the message.

```
mpirun -np 2 python pingpong.py
```

# Collective communication

- Sometimes it is convenient to share data with all the processes in the Communicator, coordinated by a root process.
- There are several collective communication operations:
  - Broadcast
  - Reduce
  - Scatter
  - Gather

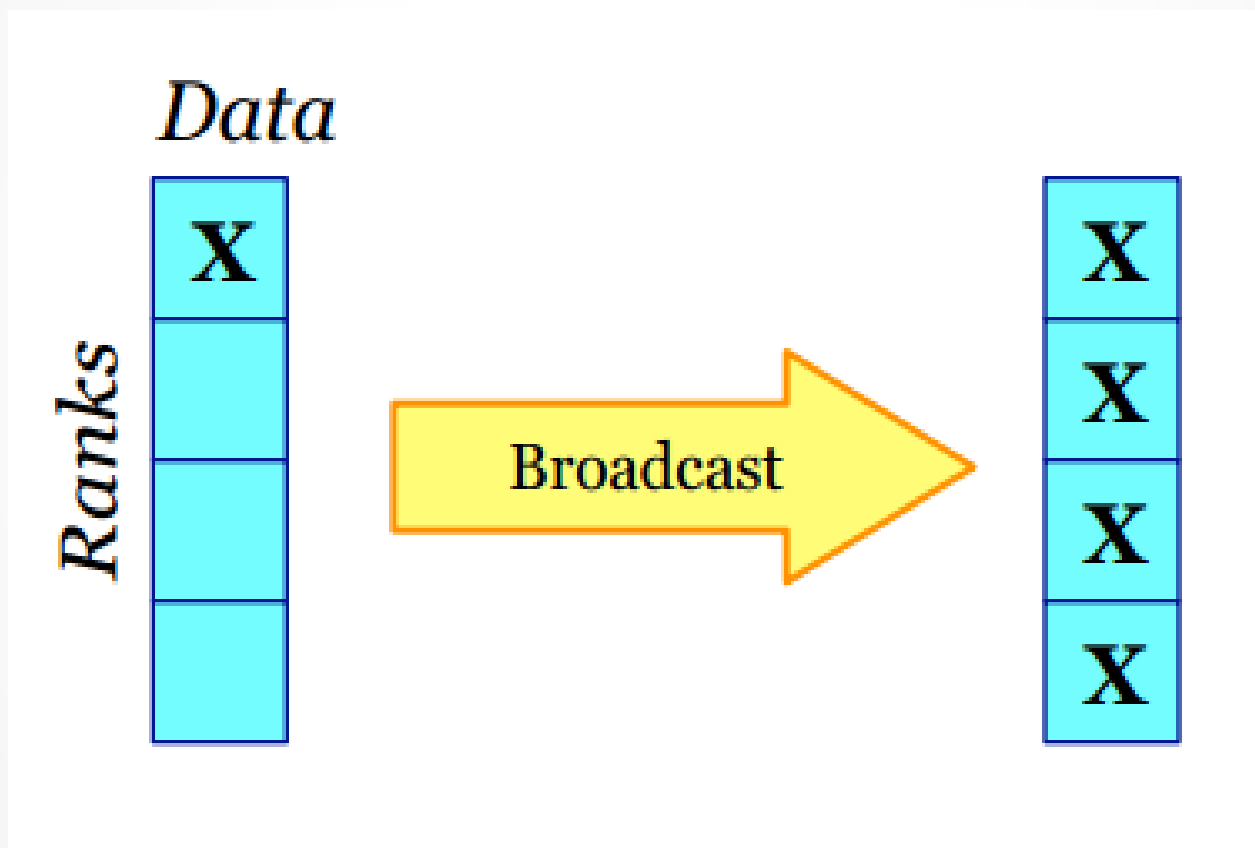
# Broadcast

- One to Many.
- A root process transmit data to all other processes.
- In the function call, the rank of the root node is specified.
- The root node transmit the data, all others receive it.



*SciPyLA*  
2016

# Broadcast




`broadcast_object.py` – `broadcast_buffer.py`



# Reduction


- Many to One.
- A mathematical operation is applied on the data in the root node indicated in the function call.
- Supported operations: Max, Min, Sum, Prod




# Scatter

- Similar to Broadcast, but instead of sending the same data to all the processes, it distributes them.

MPI\_Bcast




MPI\_Scatter






# Gather

- Opposite to scatter. Each node sends a piece of data, and all pieces are gathered in the root process.






Exercises!

# Exercise: Ring

- We'll have a ring of processes. Each process will send a message with its rank to the next. Each process will print the message it receives.
- `mpirun -np 5 python ring.py`

```
I am process 3 an received message from [2]  
I am process 4 an received message from [3]  
I am process 0 an received message from [4]  
I am process 1 an received message from [0]  
I am process 2 an received message from [1]
```



# Exercise: Distributed sum

- In the root process, create an array of 10 random integers between 1 and 100 (`numpy.random.random_integers`)
- Distribute the array among all processes as evenly as possible.
- Each node will sum the values it got (partial sum).
- Transmit the partial sum from each node to the root using Reduce.
- Sum all partial sums on the root node

# Exercise: Distributed sum

```
Data in root is: [62 44 68 60 1 70 28 53 60 29]
I am process 0 and my range goes from 0 to 2. My data: [62 44]. \ Partial sum: 106
I am process 1 and my range goes from 2 to 4. My data: [68 60]. \ Partial sum: 128
I am process 2 and my range goes from 4 to 6. My data: [ 1 70]. \ Partial sum: 71
I am process 3 and my range goes from 6 to 10. My data: [28 53 60 29]. \ Partial sum: 170
Reduced sum is: 475. From original data: 475
```

# Exercise: Parallel pi

- You are given a `pi_serial.py` script that calculates the number pi using the Wallis product:

$$\prod_{n=1}^{\infty} \left( \frac{2n}{2n-1} \cdot \frac{2n}{2n+1} \right) = \frac{2}{1} \cdot \frac{2}{3} \cdot \frac{4}{3} \cdot \frac{4}{5} \cdot \frac{6}{5} \cdot \frac{6}{7} \cdot \frac{8}{7} \cdot \frac{8}{9} \cdots = \frac{\pi}{2}$$

# Exercise: Parallel pi

- Make a parallel version, **pi\_mpi.py** that uses MPI to distribute the processing load among different processes.
- Observe the processor load while running each script (the serial and the parallel version).
- Run some benchmarks to estimate the performance enhancement.

# References

- <http://mpitutorial.com/beginner-mpi-tutorial/>
- <http://coco.sam.pitt.edu/~emeneses/teaching/mpi>
- <http://mpi4py.readthedocs.io/>